

*Annotated
Version*

Machine Learning Course - CS-433

Optimization

Sep 20+26, 2023

Martin Jaggi

Last updated on: September 20, 2023

credits to Mohammad Emtiyaz Khan



Learning / Estimation / Fitting

Given a cost function $\mathcal{L}(\mathbf{w})$, we wish to find \mathbf{w}^* which minimizes the cost:

$$\min_{\mathbf{w}} \mathcal{L}(\mathbf{w}) \quad \text{subject to } \mathbf{w} \in \mathbb{R}^D$$

This means the *learning* problem is formulated as an **optimization problem**.

We will use an **optimization algorithm** to solve the problem (to find a good \mathbf{w}).

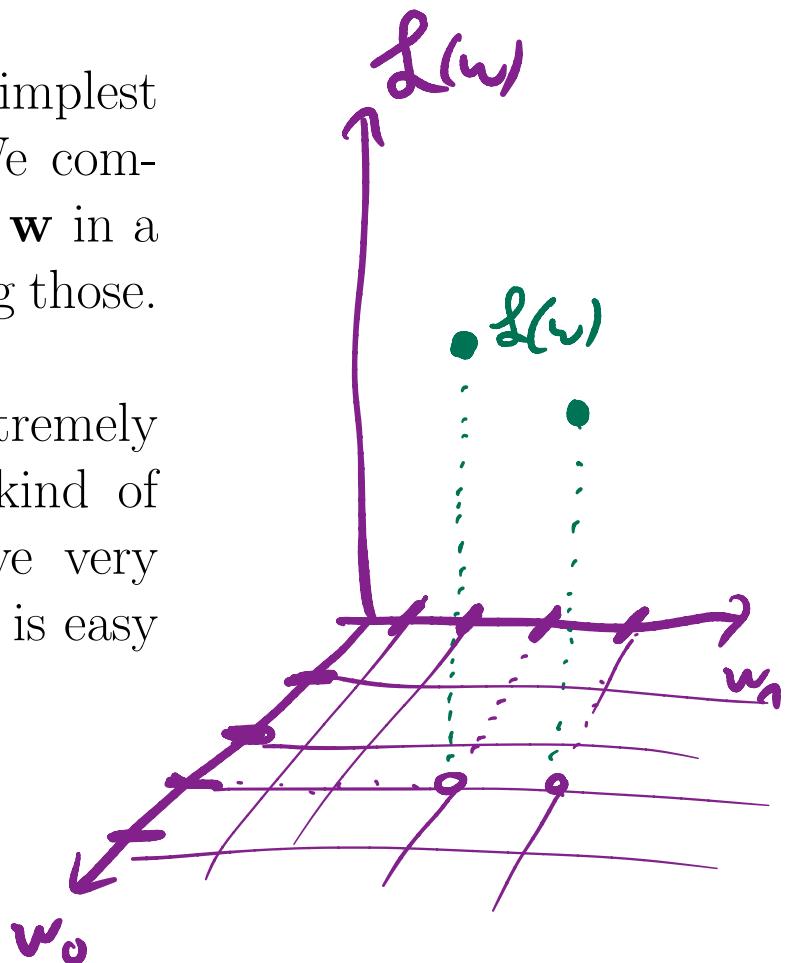
Grid Search

Grid search is one of the simplest optimization algorithms. We compute the cost over all values \mathbf{w} in a grid, and pick the best among those.

This is brute-force, but extremely simple and works for any kind of cost function when we have very few parameters and the cost is easy to compute.

$$\mathcal{L}(\mathbf{w}) = \frac{1}{N} \sum_{n=1}^N \mathcal{L}_n(\mathbf{w})$$

example: $\| \mathbf{x}_n - \mathbf{f}_{\mathbf{w}}(\mathbf{x}_n) \|_2^2$



For a large number of parameters D , however, grid search has too many “for-loops”, resulting in an exponential computational complexity:

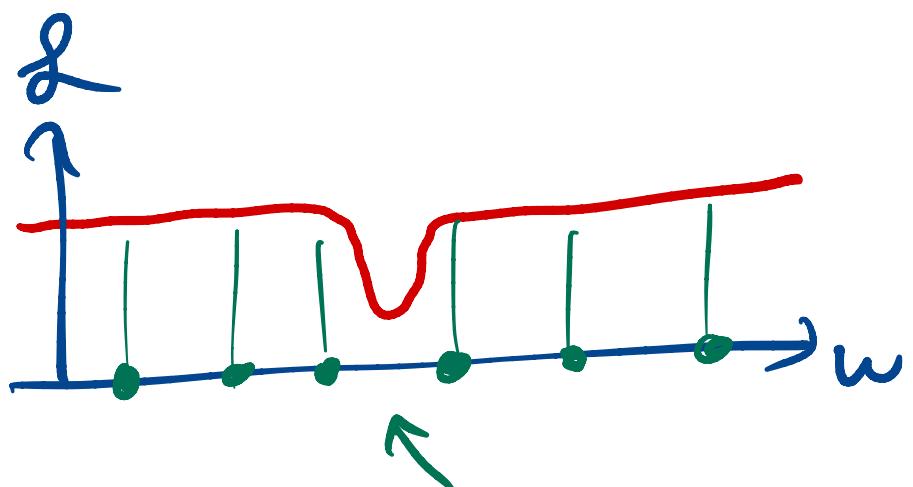
If we decide to use 10 possible values for each dimension of \mathbf{w} , then we have to check 10^D points. This is clearly impossible for most practical machine learning models, which can often have $D \approx$ millions of parameters. Choosing a good range of values for each dimension is another problem.

Other issues: No guarantee can be given that we end up close to an optimum.

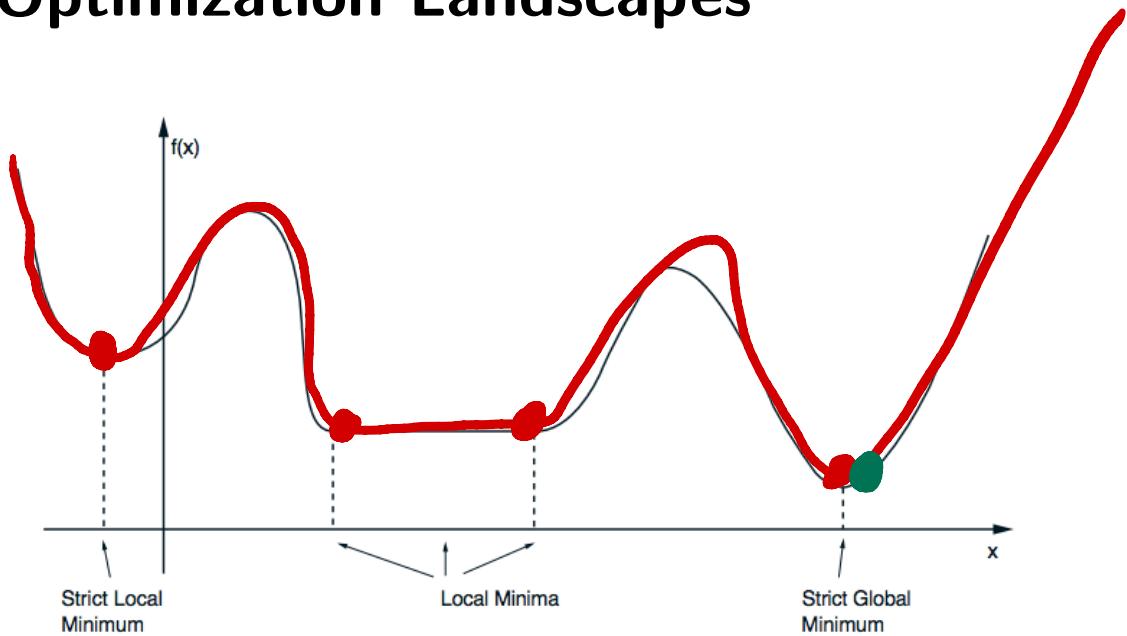
granularity

$$\begin{aligned}|\text{search range}| &= |\{1, \dots, 10\}^D| \\&= 10^D\end{aligned}$$

exponential cost



Optimization Landscapes



The above figure is taken from Bertsekas, Nonlinear programming.

A vector \mathbf{w}^* is a local minimum of \mathcal{L} if it is no worse than its neighbors; i.e. there exists an $\epsilon > 0$ such that,

$$\mathcal{L}(\mathbf{w}^*) \leq \mathcal{L}(\mathbf{w}), \quad \forall \mathbf{w} \text{ with } \|\mathbf{w} - \mathbf{w}^*\| < \epsilon$$

A vector \mathbf{w}^* is a global minimum of \mathcal{L} if it is no worse than all others,

$$\mathcal{L}(\mathbf{w}^*) \leq \mathcal{L}(\mathbf{w}), \quad \forall \mathbf{w} \in \mathbb{R}^D$$

A local or global minimum is said to be strict if the corresponding inequality is strict for $\mathbf{w} \neq \mathbf{w}^*$.

Smooth Optimization

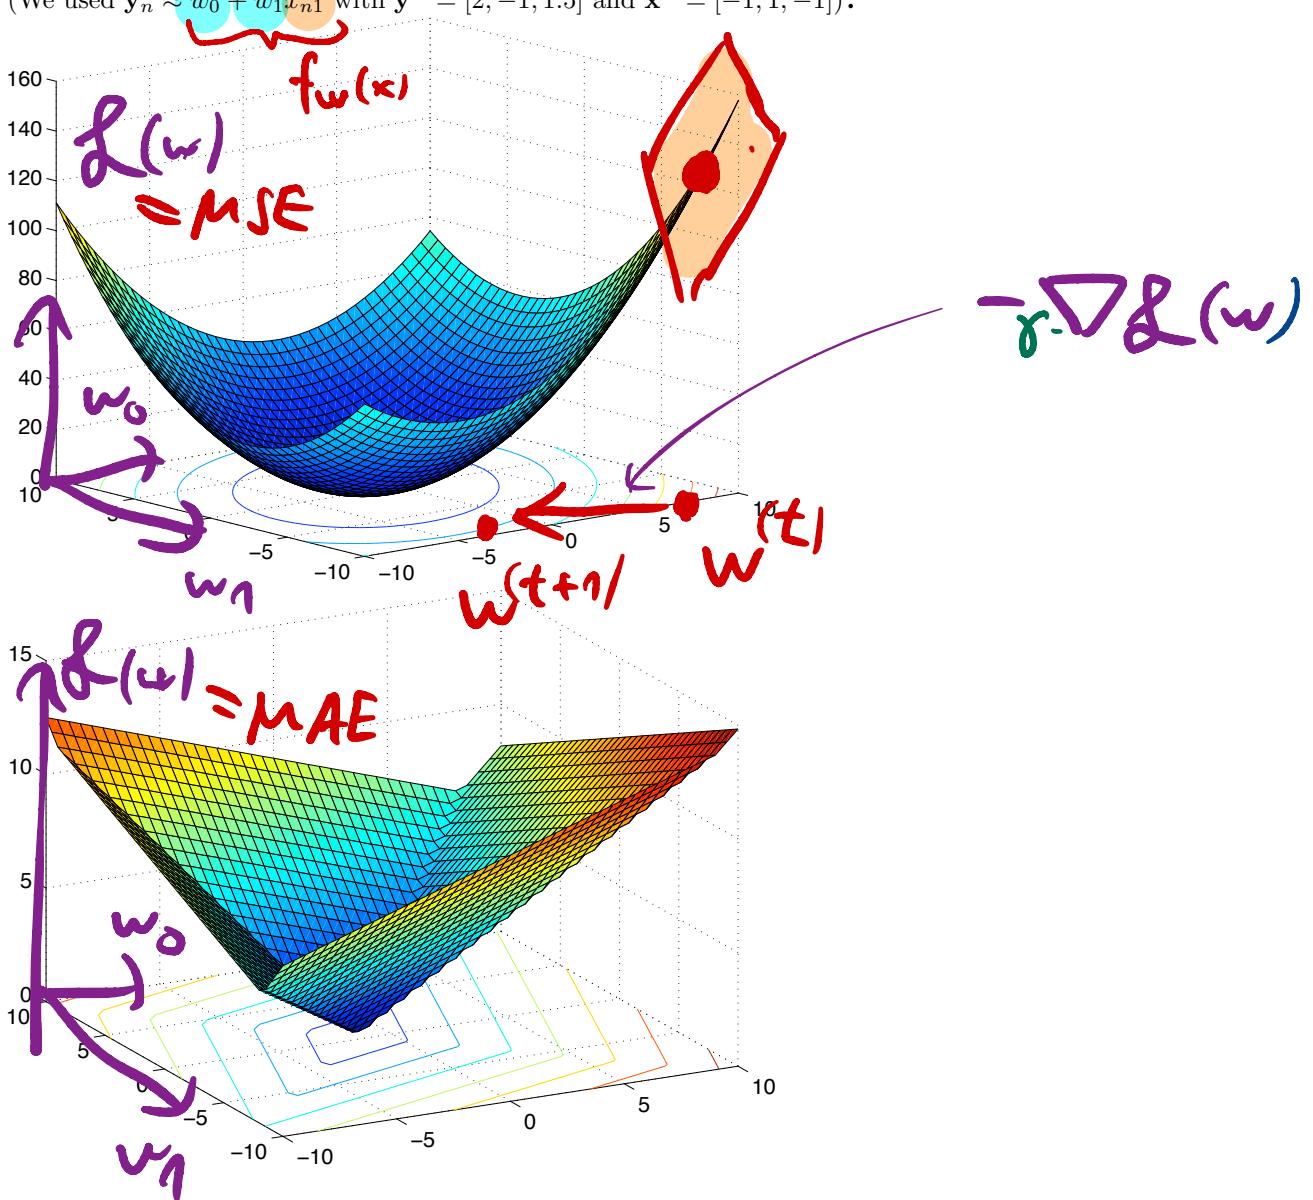
Follow the Gradient

A gradient (at a point) is the slope of the tangent to the function (at that point). It points to the direction of largest increase of the function.

Cauchy
1848

For a 2-parameter model, $\text{MSE}(\mathbf{w})$ and $\text{MAE}(\mathbf{w})$ are shown below.

(We used $\mathbf{y}_n \approx w_0 + w_1 x_{n1}$ with $\mathbf{y}^\top = [2, -1, 1.5]$ and $\mathbf{x}^\top = [-1, 1, -1]$).



Definition of the gradient:

$$\nabla \mathcal{L}(\mathbf{w}) := \left[\frac{\partial \mathcal{L}(\mathbf{w})}{\partial w_1}, \dots, \frac{\partial \mathcal{L}(\mathbf{w})}{\partial w_D} \right]^\top$$

This is a vector, $\nabla \mathcal{L}(\mathbf{w}) \in \mathbb{R}^D$.

Gradient Descent

To minimize the function, we iteratively take a step in the (opposite) direction of the gradient

$$\mathbf{w}^{(t+1)} := \mathbf{w}^{(t)} - \gamma \nabla \mathcal{L}(\mathbf{w}^{(t)})$$

where $\gamma > 0$ is the step-size (or learning rate). Then repeat with the next t .

popular stepsize
 $\gamma := \frac{\text{const}}{t}$

Example: Gradient descent for 1-parameter model to minimize MSE:

$$w_0^{(t+1)} := (1 - \gamma)w_0^{(t)} + \gamma \bar{y}$$

where $\bar{y} := \sum_n y_n / N$. When is this sequence guaranteed to converge?

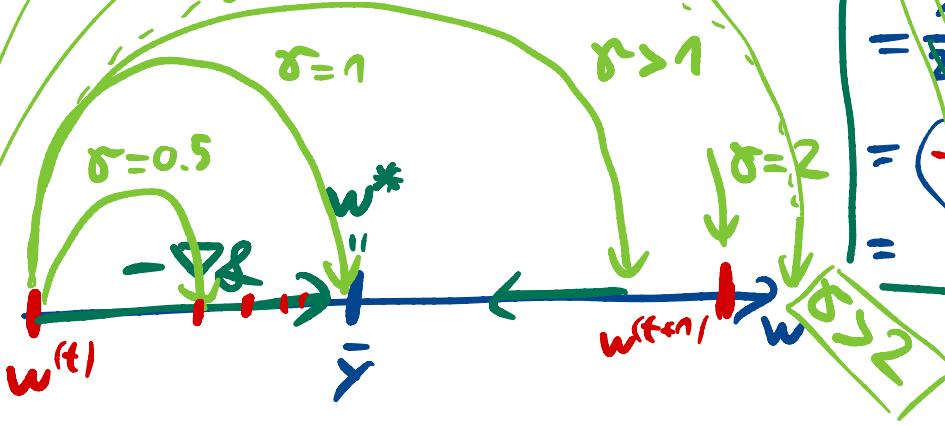
$$\min_w \mathcal{L}(w_0)$$

$$\mathcal{L}(w) = \frac{1}{2N} \sum_{n=0}^N (x_n - w_0)^2$$

$$\begin{aligned} \nabla \mathcal{L} &= \frac{\partial}{\partial w_0} \mathcal{L} \\ &= \frac{1}{2N} \sum_{n=0}^N -2(x_n - w_0) \\ &= \left(-\frac{1}{N} \sum_{n=0}^N x_n\right) + w_0 \\ &= w_0 - \bar{x} \end{aligned}$$

$\in \mathbb{R}$

divergence



Gradient Descent for Linear MSE

$$x_n \in \mathbb{R}^D$$

$$w \in \mathbb{R}^D$$

For linear regression

$$\mathbf{y} = \begin{bmatrix} y_1 \\ y_2 \\ \vdots \\ y_N \end{bmatrix}, \mathbf{X} = \underbrace{\begin{bmatrix} x_{11} & x_{12} & \dots & x_{1D} \\ x_{21} & x_{22} & \dots & x_{2D} \\ \vdots & \vdots & \ddots & \vdots \\ x_{N1} & x_{N2} & \dots & x_{ND} \end{bmatrix}}_N \quad \leftarrow \mathbf{x}_n^\top$$

We define the error vector \mathbf{e} :

$$\mathbf{e} = \mathbf{y} - \mathbf{X}\mathbf{w} \quad \leftarrow \mathbf{e} = \begin{pmatrix} e_1 \\ \vdots \\ e_N \end{pmatrix} \in \mathbb{R}^N$$

and MSE as follows:

$$\boxed{\mathcal{L}(\mathbf{w}) := \frac{1}{2N} \sum_{n=1}^N (y_n - \mathbf{x}_n^\top \mathbf{w})^2} \quad \leftarrow \mathcal{L}_n(\mathbf{w})$$

then the gradient is given by

$$\boxed{\nabla \mathcal{L}(\mathbf{w}) = -\frac{1}{N} \mathbf{X}^\top \mathbf{e}}$$

$$\begin{aligned} \frac{\partial \mathcal{L}}{\partial w_n} &= -\cancel{\mathbf{X}} \cancel{(y_n - \mathbf{x}_n^\top \mathbf{w})} \mathbf{x}_n^\top \cancel{\mathbf{w}} \\ &= -\frac{1}{N} (\mathbf{X}_{:,n})^\top \mathbf{e} \\ \frac{\partial \mathcal{L}}{\partial w_D} &= -\frac{1}{N} (\mathbf{X}_{:,D})^\top \mathbf{e} \end{aligned}$$

Computational cost. What is the complexity (# operations) of computing the gradient?

a) starting from \mathbf{w} and

b) given \mathbf{e} and \mathbf{w} ?

① compute \mathbf{e}
cost ~~$O(N \cdot D) + O(N)$~~

② compute $\nabla \mathcal{L}$, given \mathbf{e} , \mathbf{w}
 $O(N \cdot D) + O(D)$

$$\underline{O(N \cdot D)}$$

Variant with offset. Recall: Alternative trick when also incorporating an offset term for the regression:

$$\mathbf{y} = \begin{bmatrix} y_1 \\ y_2 \\ \vdots \\ y_N \end{bmatrix} \quad \tilde{\mathbf{X}} = \begin{bmatrix} 1 & x_{11} & x_{12} & \dots & x_{1D} \\ 1 & x_{21} & x_{22} & \dots & x_{2D} \\ \vdots & \vdots & \vdots & \ddots & \vdots \\ 1 & x_{N1} & x_{N2} & \dots & x_{ND} \end{bmatrix}$$

$$\mathbf{w} \in \mathbb{R}^{D+1}$$

$$\mathbf{y} - \tilde{\mathbf{X}}\mathbf{w} = \mathbf{e}$$

Artificial feature

Stochastic Gradient Descent

Sum Objectives. In machine learning, most cost functions are formulated as a sum over the training examples, that is

$$\mathcal{L}(\mathbf{w}) = \frac{1}{N} \sum_{n=1}^N \mathcal{L}_n(\mathbf{w}),$$

SGD

where \mathcal{L}_n is the cost contributed by the n -th training example.

Q: What are the \mathcal{L}_n for linear MSE?

The SGD Algorithm. The stochastic gradient descent (SGD) algorithm is given by the following update rule, at step t :

① Sample a datapoint $n \in \{1, \dots, N\}$ u.a.r.

② $\mathbf{w}^{(t+1)} := \mathbf{w}^{(t)} - \gamma \underbrace{\nabla \mathcal{L}_n(\mathbf{w}^{(t)})}_{\text{stochastic gradient}}.$

stochastic gradient = g

Theoretical Motivation. Idea:

Cheap but unbiased estimate of the gradient!

In expectation over the random choice of n , we have

$$\mathbb{E}[\nabla \mathcal{L}_n(\mathbf{w})] = \underline{\nabla \mathcal{L}(\mathbf{w})}$$

which is the true gradient direction.
(check!)

Mini-batch SGD. There is an intermediate version, using the update direction being

$$\mathbf{g} := \frac{1}{|B|} \sum_{n \in B} \nabla \mathcal{L}_n(\mathbf{w}^{(t)})$$

again with *avg over mini-batch*

$$\mathbf{w}^{(t+1)} := \mathbf{w}^{(t)} - \gamma \mathbf{g}.$$

In the above gradient computation, we have randomly chosen a subset $B \subseteq [N]$ of the training examples. For each of these selected examples n , we compute the respective gradient $\nabla \mathcal{L}_n$, at the same current point $\mathbf{w}^{(t)}$.

$$\begin{aligned} \mathbb{E}[\nabla \mathcal{L}_n(\mathbf{w})] &= \frac{1}{n} \sum_{n=1}^N \nabla \mathcal{L}_n(\mathbf{w}) \\ &= \nabla (\mathbb{E}_{\mathbf{i}}[\mathcal{L}_i(\mathbf{w})]) \\ &= \nabla \mathcal{L}(\mathbf{w}) \quad \square \end{aligned}$$

$B \subseteq \{1, \dots, N\}$
sample of size $|B|$

- $|B| = 1$
→ pure SGD
- $|B| = 128$
→ mini-batch
- $|B| = N$
→ full GD

The computation of \mathbf{g} can be parallelized easily. This is how current deep-learning applications utilize GPUs (by running over $|B|$ threads in parallel).

Note that in the extreme case $B := [N]$, we obtain (batch) gradient descent, i.e. $\mathbf{g} = \nabla \mathcal{L}$.

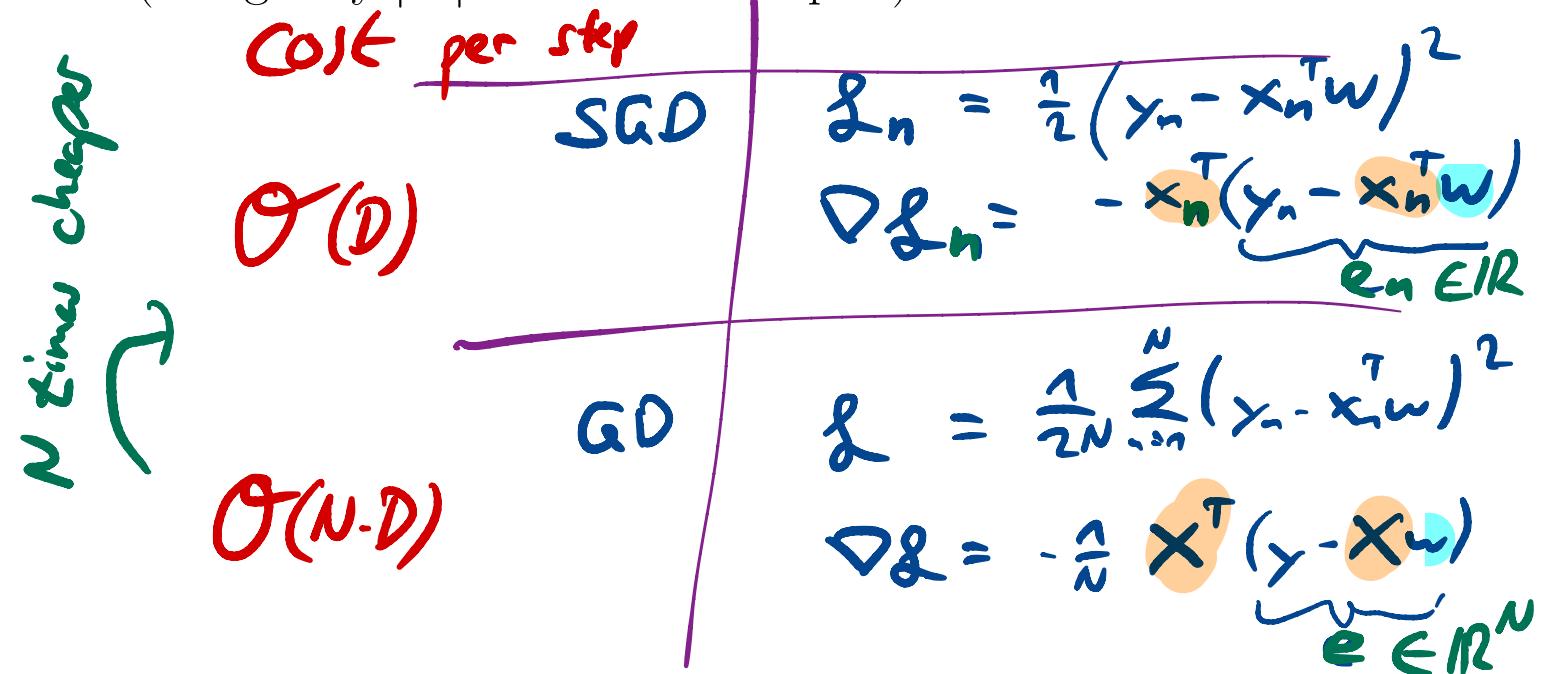
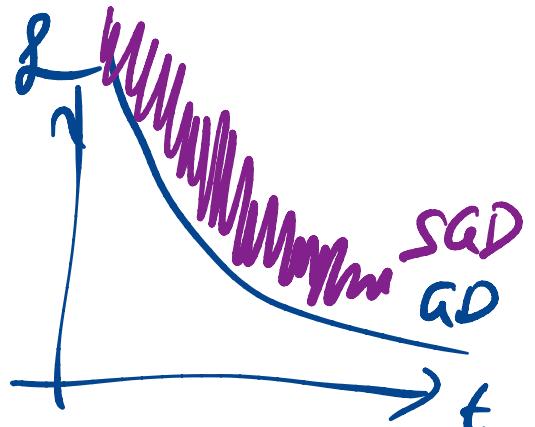
SGD for Linear MSE

See Exercise Sheet 2.



Computational cost. For linear MSE, what is the complexity (# operations) of computing the stochastic gradient?

(using only $|B| = 1$ data examples)



Non-differentiable

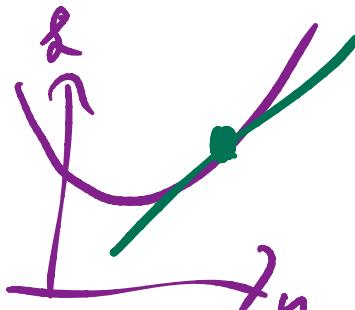
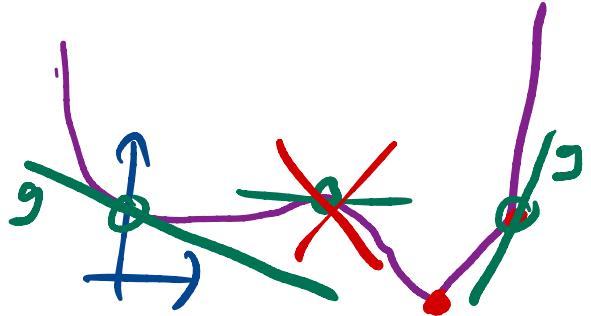
Non-Smooth Optimization

An alternative characterization of *convexity*, for differentiable functions is given by

$$\mathcal{L}(\mathbf{u}) \geq \mathcal{L}(\mathbf{w}) + \nabla \mathcal{L}(\mathbf{w})^\top (\mathbf{u} - \mathbf{w}) \quad \forall \mathbf{u}, \mathbf{w}$$

1st-order Taylor exp.

meaning that the function must always lie above its *linearization*.

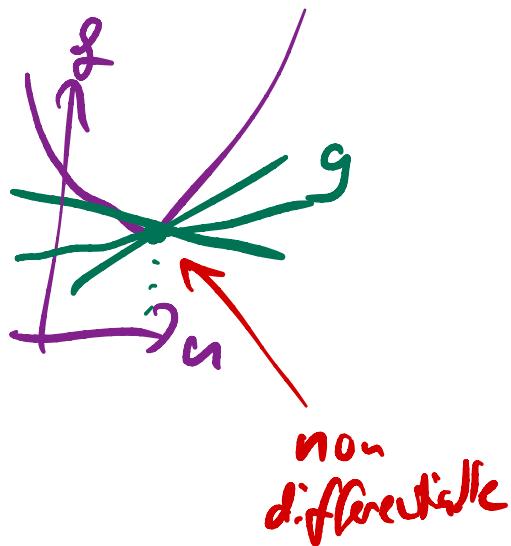


Subgradients

A vector $\mathbf{g} \in \mathbb{R}^D$ such that

$$\mathcal{L}(\mathbf{u}) \geq \mathcal{L}(\mathbf{w}) + \mathbf{g}^\top (\mathbf{u} - \mathbf{w}) \quad \forall \mathbf{u}$$

is called a *subgradient* to the function \mathcal{L} at \mathbf{w} .



This definition makes sense for objectives \mathcal{L} which are not necessarily differentiable (and not even necessarily convex).

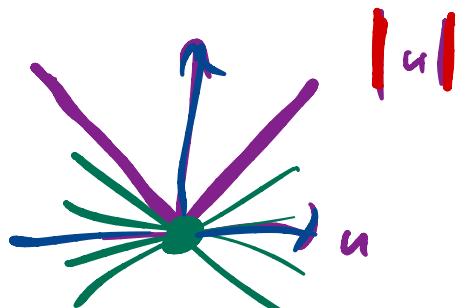
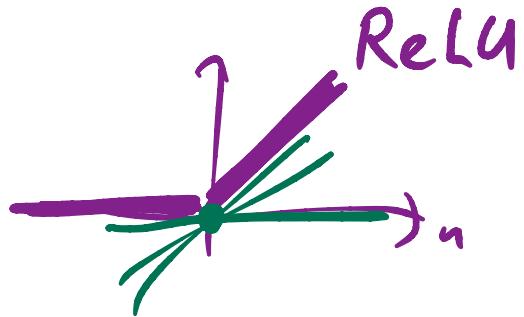
If \mathcal{L} is convex and differentiable at \mathbf{w} , then the only subgradient at \mathbf{w} is $\mathbf{g} = \nabla \mathcal{L}(\mathbf{w})$.

Subgradient Descent

Identical to the gradient descent algorithm, but using a subgradient instead of gradient. Update rule

$$\mathbf{w}^{(t+1)} := \mathbf{w}^{(t)} - \gamma \mathbf{g}$$

for \mathbf{g} being a subgradient to \mathcal{L} at the current iterate $\mathbf{w}^{(t)}$.



Example: Optimizing Linear MAE

1. Compute a subgradient of the absolute value function

$$h : \mathbb{R} \rightarrow \mathbb{R}, h(e) := |e|.$$

$$g = \begin{cases} -1 & \text{if } e < 0 \\ [-1, 1] & \text{if } e = 0 \\ 1 & \text{if } e > 0 \end{cases}$$

2. Recall the definition of the mean absolute error:

$$\mathcal{L}(\mathbf{w}) = \text{MAE}(\mathbf{w}) := \frac{1}{N} \sum_{n=1}^N |y_n - f_{\mathbf{w}}(\mathbf{x}_n)|$$

For linear regression, its (sub)gradient is easy to compute using the chain rule.
Compute it!

$$\mathbf{x}_n^T \mathbf{w}$$

$$||$$

$$f(w) = h(q(w))$$

h : non-differentiable
 q : differentiable

Subgradient of f at w

See Exercise Sheet 2.

$$g \in \boxed{\partial h(q(w))} \cdot \nabla q(w)$$

$\partial h = \text{set of subgradients of } h$

Stochastic (Sub)gradient Descent

Stochastic SubGradient Descent
(still abbreviated SGD commonly).

Same, \mathbf{g} being a subgradient to the randomly selected \mathcal{L}_n at the current iterate $\mathbf{w}^{(t)}$.

$$\mathcal{L}(\mathbf{w}) = \frac{1}{N} \sum_{n=1}^N \mathcal{L}_n(\mathbf{w})$$

↑ ↑
potentially
not differentiable

Exercise: Compute the SGD update for linear MAE.

$\mathbf{g} :$

take a subgradient
of \mathcal{L}_n

Linear model training

	smooth MSE	non-smooth MAE
GD	gradient $\nabla \mathcal{L}$ $\mathcal{O}(N \cdot D)$	subgradient $\mathbf{g} \in \partial \mathcal{L}$ $\mathcal{O}(N \cdot D)$
SGD	gradient $\nabla \mathcal{L}_n$ $\mathcal{O}(D)$	subgradient $\mathbf{g} \in \partial \mathcal{L}_n$ $\mathcal{O}(D)$
per iteration		

Variants of SGD

SGD with Momentum

$$\beta_1 = 0.9$$

pick a stochastic gradient \mathbf{g}

$$\mathbf{m}^{(t+1)} := \beta_1 \mathbf{m}^{(t)} + (1 - \beta_1) \mathbf{g} \quad (\text{momentum term})$$

$$\mathbf{w}^{(t+1)} := \mathbf{w}^{(t)} - \gamma \mathbf{m}^{(t+1)}$$

g in standard SGD

- momentum from previous gradients (acceleration)

Adam

pick a stochastic gradient \mathbf{g}

$$\mathbf{m}^{(t+1)} := \beta_1 \mathbf{m}^{(t)} + (1 - \beta_1) \mathbf{g} \quad (\text{momentum term})$$

$$\mathbf{v}_i^{(t+1)} := \beta_2 \mathbf{v}_i^{(t)} + (1 - \beta_2) (\mathbf{g}_i)^2 \quad \forall i \quad (\text{2nd-order statistics})$$

$$\mathbf{w}_i^{(t+1)} := \mathbf{w}_i^{(t)} - \frac{\gamma}{\sqrt{\mathbf{v}_i^{(t+1)}}} \mathbf{m}_i^{(t+1)} \quad \forall i$$

- faster forgetting of older weights
- is a momentum variant of Adagrad
- momentum from previous gradients (acceleration)
- strong performance in practice, e.g. for self-attention networks

$$\frac{\mathbf{m}}{\sqrt{\mathbf{v}}} = \frac{\mathbf{g}}{\sqrt{\mathbf{g}^2}} = \frac{\mathbf{g}}{|\mathbf{g}|} = \text{sign}(\mathbf{g})$$

SignSGD

pick a stochastic gradient \mathbf{g}

$$\mathbf{w}_i^{(t+1)} := \mathbf{x}_i^{(t)} - \gamma \text{sign}(\mathbf{g}_i) \quad \forall i$$

- only use the sign (one bit) of each gradient entry - efficient for distributed training
- convergence issues

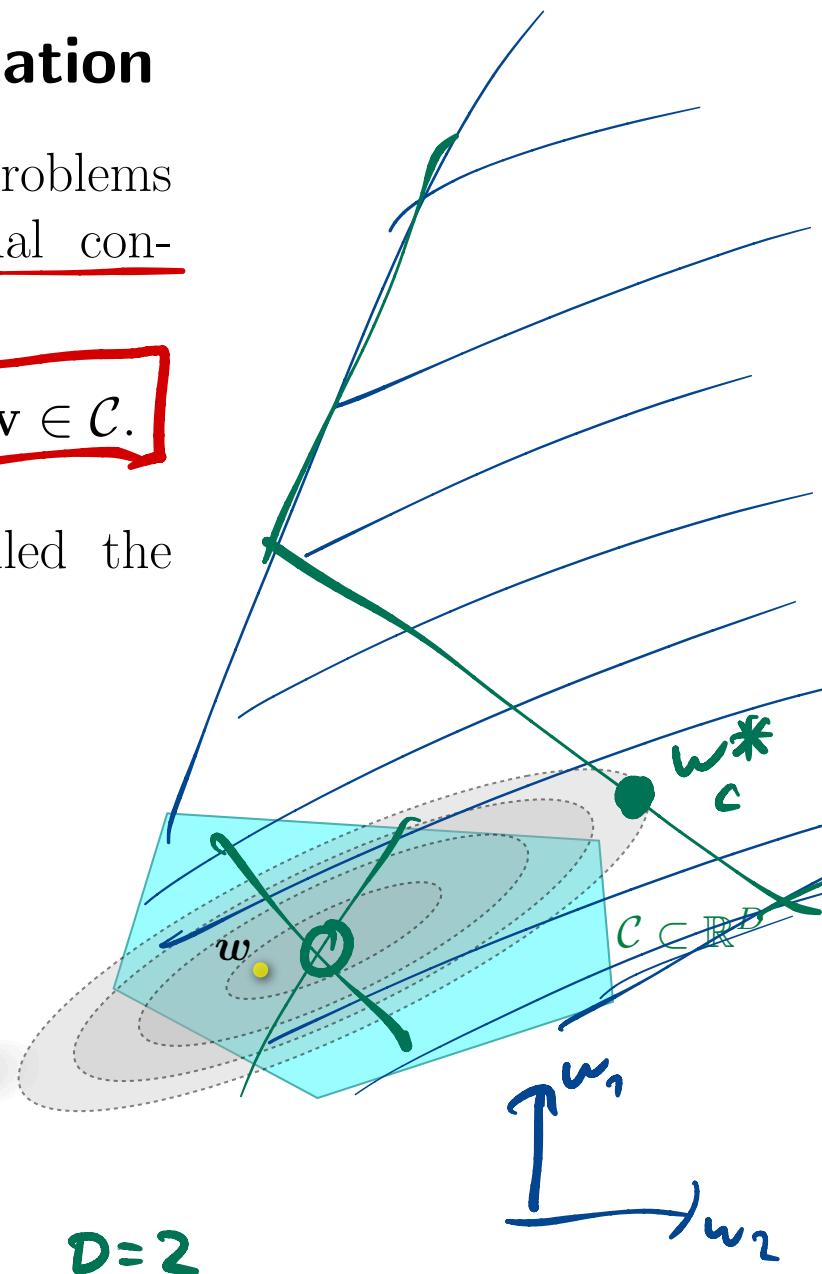
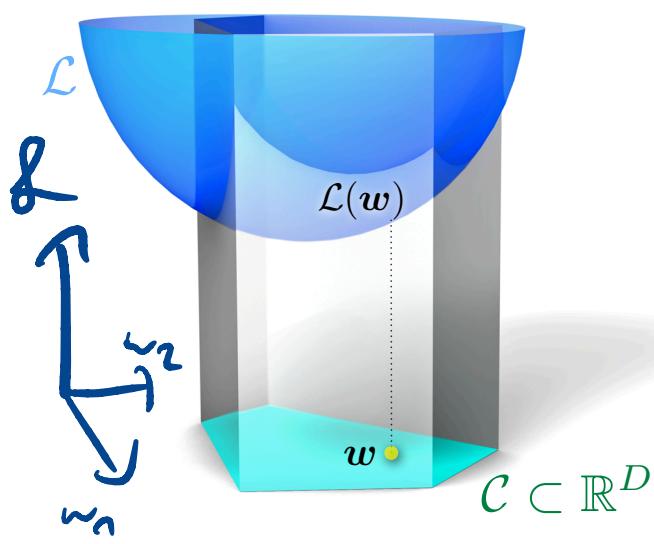
special case of Adam with $\beta_1 = 0$ and $\beta_2 = 0$

Constrained Optimization

Sometimes, optimization problems come posed with additional constraints:

$$\min_{\mathbf{w}} \mathcal{L}(\mathbf{w}), \quad \text{subject to } \mathbf{w} \in \mathcal{C}.$$

The set $\mathcal{C} \subset \mathbb{R}^D$ is called the constraint set.



Solving Constrained Optimization Problems

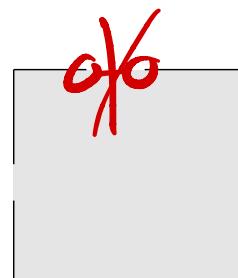
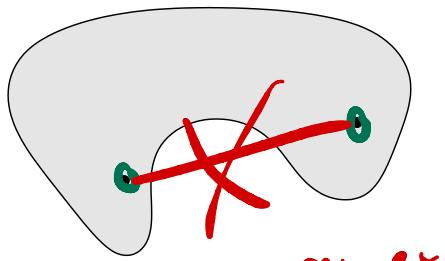
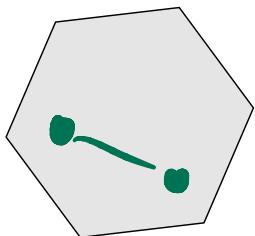
- A) Projected Gradient Descent
- B) Transform it into an *unconstrained* problem

Convex Sets

A set \mathcal{C} is **convex** iff

the line segment between any two points of \mathcal{C} lies in \mathcal{C} , i.e., if for any $\mathbf{u}, \mathbf{v} \in \mathcal{C}$ and any θ with $0 \leq \theta \leq 1$, we have

$$\theta\mathbf{u} + (1 - \theta)\mathbf{v} \in \mathcal{C}.$$



non-convex

*Figure 2.2 from S. Boyd, L. Vandenberghe

Properties of Convex Sets

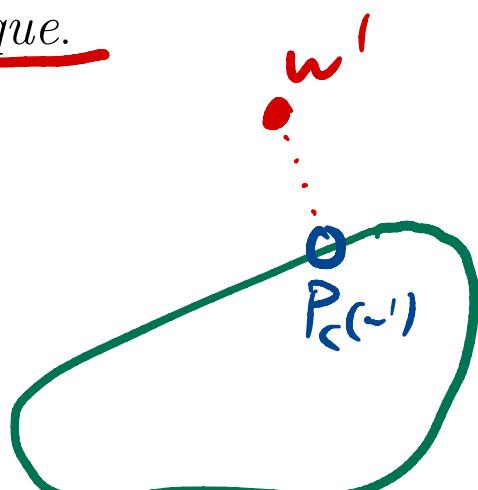
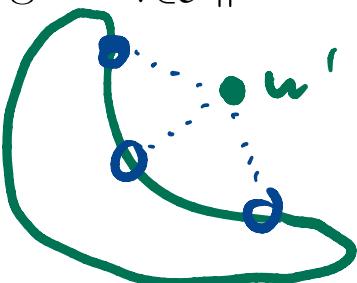
- Intersections of convex sets are convex

- Projections onto convex sets are unique.
(and often efficient to compute)

Formal definition:

$$P_{\mathcal{C}}(\mathbf{w}') := \arg \min_{\mathbf{v} \in \mathcal{C}} \|\mathbf{v} - \mathbf{w}'\|^2$$

non-convex



Projected Gradient Descent

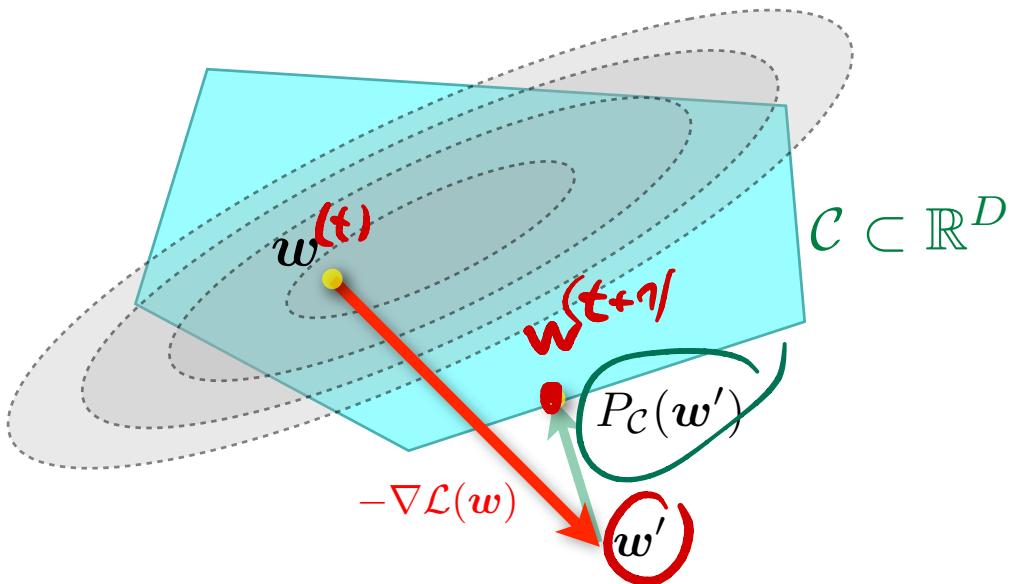
Idea: add a projection onto \mathcal{C} after every step:

$$P_{\mathcal{C}}(\mathbf{w}') := \arg \min_{\mathbf{v} \in \mathcal{C}} \|\mathbf{v} - \mathbf{w}'\|.$$

Update rule:

$$\mathbf{w}^{(t+1)} := P_{\mathcal{C}} \left[\mathbf{w}^{(t)} - \gamma \nabla \mathcal{L}(\mathbf{w}^{(t)}) \right].$$

traditional GD step = \mathbf{w}'



Projected SGD. Same SGD step followed by the projection step, as above. Same convergence properties.

Computational cost of projection?
Crucial!

δ

$$\min_{\mathbf{w} \in \mathcal{C}} \delta(\mathbf{w})$$

w^*

$$\min_{\mathbf{w}} \delta(\mathbf{w})$$

w^*_c

2

$$\min_{w \in C} L(w) \rightsquigarrow \min_w L(w) + P(w)$$

Turning Constrained into Unconstrained Problems

(Alternatives to projected gradient methods)

Use penalty functions instead of directly solving $\min_{w \in C} \mathcal{L}(w)$.

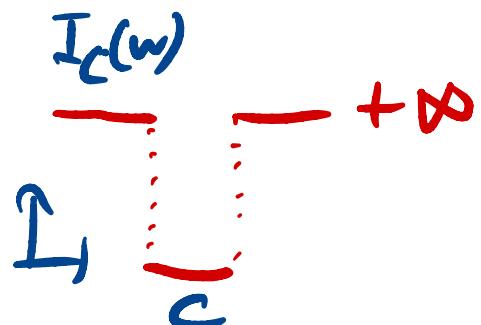
Penalty function

- “brick wall” (indicator function)

$$P(w) = I_C(w) := \begin{cases} 0 & w \in C \\ \infty & w \notin C \end{cases}$$

$$\Rightarrow \min_{w \in \mathbb{R}^D} \mathcal{L}(w) + I_C(w)$$

(disadvantage: non-continuous objective)

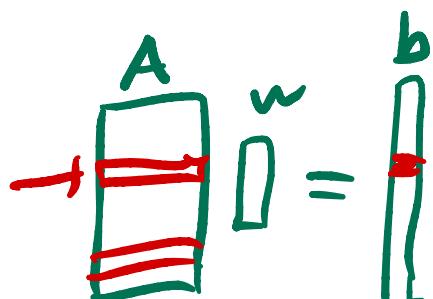


- Penalize error. Example:

$$C = \{w \in \mathbb{R}^D \mid Aw = b\}$$

$$\Rightarrow \min_{w \in \mathbb{R}^D} \mathcal{L}(w) + \lambda \|Aw - b\|^2$$

“distance” from the set C



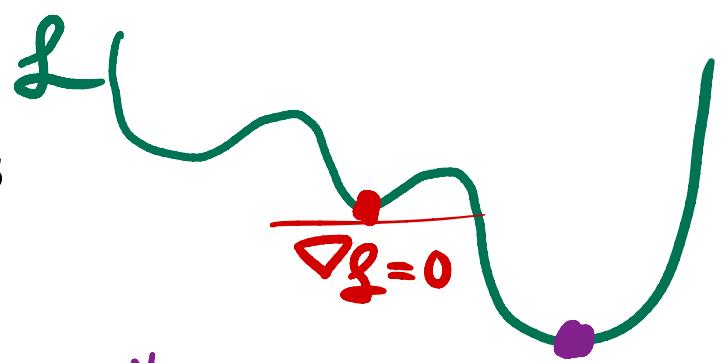
- Linearized Penalty Functions (see Lagrange Multipliers)

• Sparsity

$$P(w) := \lambda \|w\|_1$$

Implementation Issues

For gradient methods:



Stopping criteria: When $\|\nabla L(w)\|$ is (close to) zero, we are (often) close to ~~the~~ optimum value. (*local opt*)

Optimality: If the second-order derivative is positive (positive semi-definite to be precise), then it is a (possibly local) minimum. If the function is also convex, then this condition implies that we are at a global optimum. See the supplementary section on Optimality Conditions.

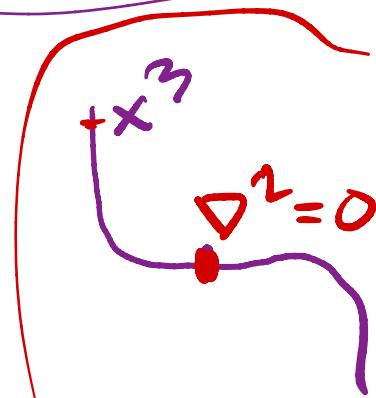
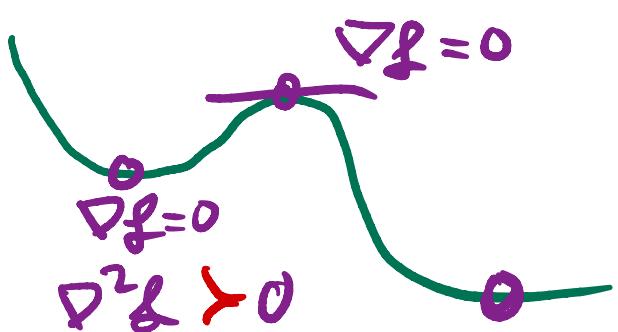
1-st order optimality condition

if ① \mathcal{L} convex
+ ② $\nabla \mathcal{L}(w) = 0$
 \Rightarrow global optimality

2nd - order optimality condition

Let \mathcal{L} be potentially non-convex
and ① $\nabla \mathcal{L}(w) = 0$
+ ② $\nabla^2 \mathcal{L}(w) > 0$

$\Rightarrow w$ is a local optimum

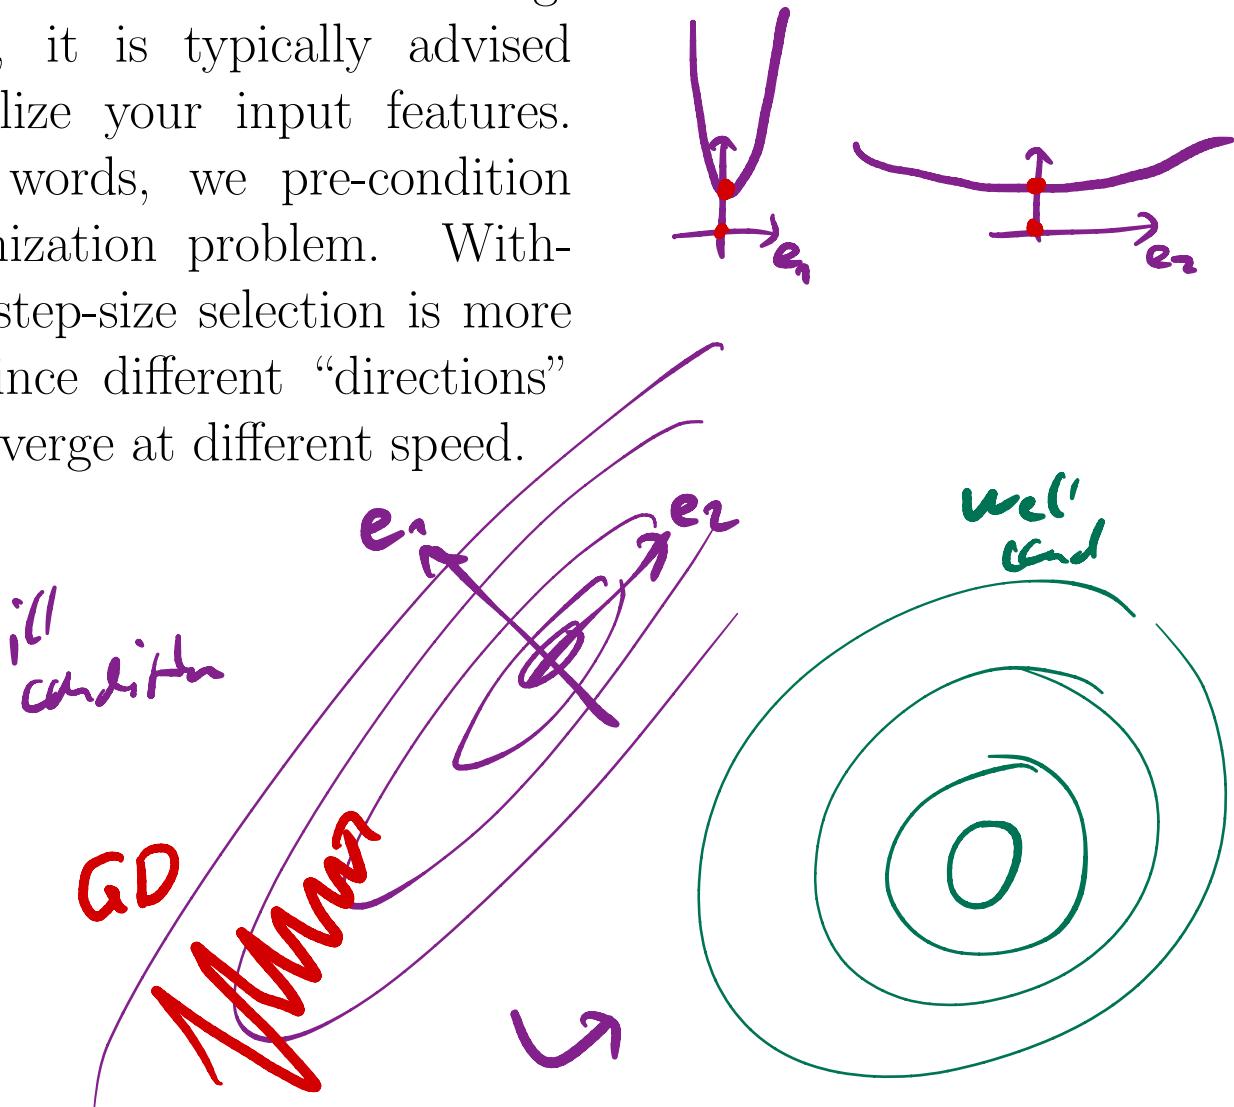


Step-size selection: If γ is too big, the method might diverge. If it is too small, convergence is slow. Convergence to a local minimum is guaranteed only when $\gamma < \gamma_{min}$ where γ_{min} is a fixed constant that depends on the problem.

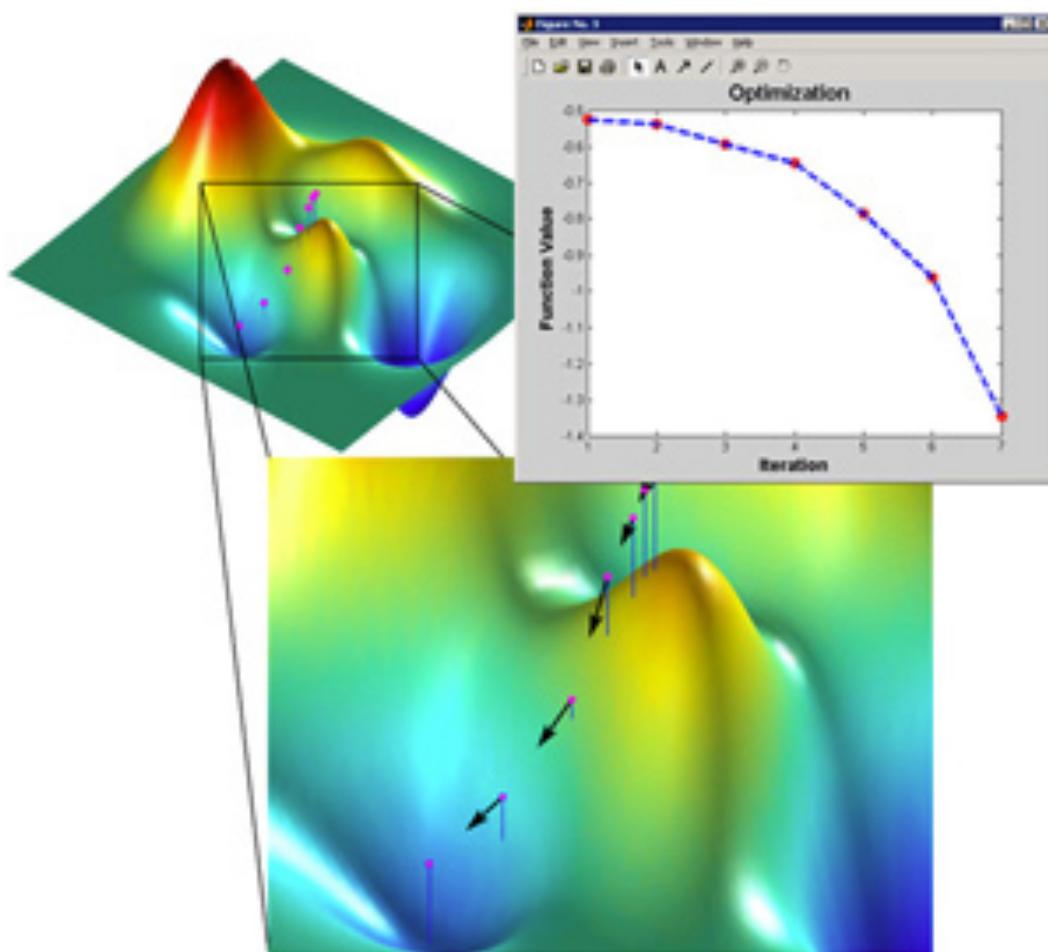
Adam is using
coordinate-wise
learning-rate

Line-search methods: For some objectives \mathcal{L} , we can set step-size automatically using a line-search method. More details on “backtracking” methods can be found in Chapter 1 of Bertsekas’ book on “nonlinear programming”.

Feature normalization and pre-conditioning: Gradient descent is very sensitive to ill-conditioning. Therefore, it is typically advised to normalize your input features. In other words, we pre-condition the optimization problem. Without this, step-size selection is more difficult since different “directions” might converge at different speed.



Non-Convex Optimization



*image from mathworks.com

Real-world problems are **not convex!**

All we have learnt on algorithm design and performance of convex algorithms still helps us in the non-convex world.

Additional Notes

Grid Search and Hyper-Parameter Optimization

Read more about grid search and other methods for “hyperparameter” setting:

en.wikipedia.org/wiki/Hyperparameter_optimization#Grid_search.

Computational Complexity

The computation cost is expressed using the [big-O](#) notation. Here is a definition taken from Wikipedia. Let f and g be two functions defined on some subset of the real numbers. We write $f(x) = \mathcal{O}(g(x))$ as $x \rightarrow \infty$, if and only if there exists a positive real number c and a real number x_0 such that $|f(x)| \leq c|g(x)|$, $\forall x > x_0$.

Please read and learn more from this page in Wikipedia:

en.wikipedia.org/wiki/Computational_complexity_of_mathematical_operations#Matrix_algebra .

- What is the computational complexity of matrix multiplication?
- What is the computational complexity of matrix-vector multiplication?

Optimality Conditions

For a *convex* optimization problem, the first-order *necessary* condition says that at *an* optimum the gradient is equal to zero.

$$\nabla \mathcal{L}(\mathbf{w}^*) = \mathbf{0} \quad (1)$$

The second-order *sufficient* condition ensures that the optimum is a minimum (not a maximum or saddle-point) using the [Hessian](#) matrix,

which is the matrix of second derivatives:

$$\mathbf{H}(\mathbf{w}^*) := \frac{\partial^2 \mathcal{L}(\mathbf{w}^*)}{\partial \mathbf{w} \partial \mathbf{w}^\top} \quad \text{is positive semi-definite.} \quad (2)$$

The Hessian is also related to the convexity of a function: a twice-differentiable function is convex if and only if the Hessian is positive semi-definite at all points.

SGD Theory

As we have seen above, when N is large, choosing a random training example (\mathbf{x}_n, y_n) and taking an SGD step is advantageous:

$$\mathbf{w}^{(t+1)} := \mathbf{w}^{(t)} - \gamma^{(t)} \nabla \mathcal{L}_n(\mathbf{w}^{(t)})$$

For convergence, $\gamma^{(t)} \rightarrow 0$ “appropriately”. One such condition called the Robbins-Monroe condition suggests to take $\gamma^{(t)}$ such that:

$$\sum_{t=1}^{\infty} \gamma^{(t)} = \infty, \quad \sum_{t=1}^{\infty} (\gamma^{(t)})^2 < \infty \quad (3)$$

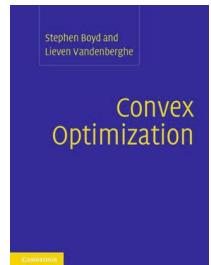
One way to obtain such sequences is $\gamma^{(t)} := 1/(t+1)^r$ where $r \in (0.5, 1)$.

More Optimization Theory

If you want, you can gain a deeper understanding of several optimization methods relevant for machine learning from this survey:

Convex Optimization: Algorithms and Complexity

- by Sébastien Bubeck



And also from the book of Boyd & Vandenberghe
(both are free online PDFs)

Exercises

1. Chain-rule



If it has been a while, familiarize yourself with it again.

2. Revise computational complexity (also see the Wikipedia link in Page 6 of lecture notes).
3. Derive the computational complexity of grid-search, gradient descent and stochastic gradient descent for linear MSE (# steps and cost per step).
4. Derive the gradients for the linear MSE and MAE cost functions.
5. Implement gradient descent and gain experience in setting the step-size.
6. Implement SGD and gain experience in setting the step-size.